

Automated Image Tagging System

Project Documentation

Team members

Ahmed Waheed: Data Engineer

Bedour Fouad: Data Engineer

Menna Osama: Model Development

Norhan El-Sayed: Model Development

Mina Farid: Deployment and Documentation

Project Overview

The objective was to create a machine learning application capable of performing multi-label classification for image tagging, leveraging a subset of the COCO dataset and AWS services to train and deploy the model.

Key Features

- **Dataset:** Subset of COCO dataset, optimized for multi-label classification with 80 classes and 100,000 training examples.
 - **Data Processing:** Data integrity checks, preprocessing, and augmentation.
 - **Model Architecture:** Transfer learning using ResNet50.
 - **Training:** Conducted on AWS SageMaker using a SageMaker notebook.
 - **Deployment:** Model deployed on AWS using SageMaker Inference Endpoints and API Gateway.
 - **Monitoring and Debugging:** AWS CloudWatch Logs were utilized to debug and monitor both the SageMaker endpoint and API Gateway.
 - **App Integration:** Simple Flutter application to demo the model via API calls.
-

Dataset

We selected a subset of the COCO dataset that is specifically optimized for multi-label classification. The dataset contains:

- **80 classes:** Various object categories relevant to the COCO dataset.
 - **100,000 training examples:** Diverse images representing the classes. The dataset was uploaded and stored in AWS S3.
-

Data Integrity Checks

Before training the model, we conducted a thorough review of the dataset to ensure the following:

- **Consistency of labels:** We ensured that the labels are correct and represent the images accurately.
 - **Completeness of data:** We ensured that the images had corresponding labels and no missing data.
 - **Data distribution:** The dataset was checked for class imbalance to ensure diversity in training.
-

Data Preprocessing and Augmentation

Data preprocessing was done to standardize image sizes and enhance model performance. Techniques included:

- **Resizing images** to match the input size for the ResNet50 model.
 - **Normalization** to ensure pixel values are between 0 and 1.
 - **Data Augmentation:** Applied techniques like horizontal flipping, random cropping, and rotations to increase the diversity of the training set.
-

Model Development

We used **transfer learning** with a **ResNet50** pre-trained model. The rationale for selecting ResNet50 was its effectiveness in handling complex image classification tasks while allowing the model to generalize well for our multi-label classification problem.

Training

Training was carried out on **AWS SageMaker** using a SageMaker notebook instance. Key details:

- **Pre-trained model:** Transfer learning with ResNet50.
- **Validation:** The model's performance was validated using a separate test set.

During the training process, **AWS CloudWatch Logs** were used to monitor the model's performance, track potential errors, and review training outputs. The logs provided valuable debugging information that helped fine-tune the training pipeline. After training, the model was saved and stored in AWS S3 for future access.

Deployment

We deployed the trained model using **AWS SageMaker Serverless Inference Endpoints**. To facilitate interaction with the model, we integrated:

- **API Gateway:** Set up to create RESTful API endpoints for the model.
- **API Keys:** Generated to control access to the model and enable authentication for requests.

AWS CloudWatch Logs were crucial for debugging and monitoring both the **SageMaker Endpoint** and **API Gateway**.

Specifically, the logs helped us:

- Debug API errors and failures, such as request failures and endpoint timeouts.
 - Monitor the **number of inference requests** sent to the API Gateway, ensuring smooth operation and tracking usage. This setup allows for easy scalability and access to the model for inference purposes.
-

Application Development

To demonstrate the use of the trained model, we built a simple **Flutter application**. The app:

- **Connects to the API:** Uses API calls to communicate with the deployed model on AWS.
 - **Image input:** Allows users to upload images for tagging and classification.
 - **Displays results:** The model returns the top predicted labels, which are displayed on the app interface
-

Conclusion

This project successfully developed a machine learning solution for **multi-label image classification** using a subset of the COCO dataset, **ResNet50** for transfer learning, and **AWS SageMaker** for training and deployment. We ensured data integrity, applied preprocessing and augmentation, and deployed the model with **SageMaker Inference Endpoints** and **API Gateway**.

A **Flutter app** was created to demonstrate the model's real-time tagging capabilities. Throughout the process, **AWS CloudWatch Logs** were used for debugging and monitoring API requests, ensuring smooth model access and performance. The project highlights the scalability and effectiveness of integrating cloud-based infrastructure for machine learning solutions.

Additional Resources

Demo App

You can try the demo app online using the following links:

[Demo App \(APK\)](#)

[Demo App \(Windows\)](#)

GitHub Repository

The full source code and additional information can be found in the GitHub repository:

[GitHub Repository](#)